

AgenticCyOps

Revision results report

9788 trials (8283 attack, 1505 benign, 0 errors) in 4 primary-model groups: llama8b_div4, mistral_div3p, q235_div4, scout_div4

Main group q235_div4: 4 domains, 3 configurations, 7 ablation configurations, persistent-state benign runs, 4 ASB validator panels

Scoring v3 effects oracle; Wilson and cluster-bootstrap intervals; Holm-corrected paired differences

Code: git 7ec5bf2 | defense freeze: defense-freeze-v2.1


Generated 2026-09-21 01:36 from results/paper_tables.md

Models and validator panels per group

Group	Primary model	Quant	Panel	Validators	vLLM	Freeze tag
llama8b_div4	meta-llama/Llama-3.1-8B-Instruct	bf16	div4	V1 Qwen3-32B; V5 Mistral-Small-3.2-24B-Instruct-2506; V4 claude-sonnet-4-5-20250929; V6 gpt-4o (threshold 3)	0.21.0	defense-freeze-v2-1-gbe6b4f1
llama8b_div4	meta-llama/Llama-3.1-8B-Instruct	bf16	-	-	0.21.0	defense-freeze-v2-1-gbe6b4f1
mistral_div3p	mistralai/Mistral-Small-3.2-24B-Instruct-2506	bf16	mistral_div3p	V1 Qwen3-32B; V4 claude-sonnet-4-5-20250929; V6 gpt-4o (threshold 2)	0.19.1rc1.dev117+g3352bf8b0	defense-freeze-v2-8-g98dde98
mistral_div3p	mistralai/Mistral-Small-3.2-24B-Instruct-2506	bf16	-	-	0.19.1rc1.dev117+g3352bf8b0	defense-freeze-v2-8-g98dde98
q235_div4	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	V1 Qwen3-32B; V5 Mistral-Small-3.2-24B-Instruct-2506; V4 claude-sonnet-4-5-20250929; V6 gpt-4o (threshold 3)	0.19.1rc1.dev117+g3352bf8b0	defense-freeze-v2-1-gbe6b4f1
q235_div4	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	-	-	0.19.1rc1.dev117+g3352bf8b0	defense-freeze-v2-1-gbe6b4f1
scout_div4	meta-llama/Llama-4-Scout-17B-16E-Instruct	bf16	div4	V1 Qwen3-32B; V5 Mistral-Small-3.2-24B-Instruct-2506; V4 claude-sonnet-4-5-20250929; V6 gpt-4o (threshold 3)	0.19.1rc1.dev117+g3352bf8b0	defense-freeze-v2-8-g98dde98
scout_div4	meta-llama/Llama-4-Scout-17B-16E-Instruct	bf16	-	-	0.19.1rc1.dev117+g3352bf8b0	defense-freeze-v2-8-g98dde98

One row per (group, primary, panel) seen in runs.csv; flat and ACL-hardened runs have no panel.

T1. Headline attack success by primary model group




T1. Headline attack success by group and configuration


Group	Config	N	ASR % [Wilson 95%]	Cluster bootstrap 95%	Attempt %	Block given attempt %	Δ vs AgenticCyOps
llama8b_div4	Flat	378	37.8 [33.1, 42.8]	[29.9, 45.8]	37.8	0.0	35.4 pp [27.5, 43.4], p=0.000
llama8b_div4	ACL-Hardened	378	37.0 [32.3, 42.0]	[29.1, 45.0]	37.6	1.4	34.7 pp [26.7, 42.6], p=0.000
llama8b_div4	AgenticCyOps	378	2.4 [1.3, 4.5]	[0.3, 5.0]	40.7	94.2	—
mistral_div3p	Flat	378	23.5 [19.6, 28.1]	[16.9, 30.7]	23.5	0.0	22.0 pp [15.3, 29.1], p=0.000
mistral_div3p	ACL-Hardened	378	23.0 [19.1, 27.5]	[16.4, 30.2]	24.3	5.4	21.4 pp [14.8, 28.6], p=0.000
mistral_div3p	AgenticCyOps	378	1.6 [0.7, 3.4]	[0.0, 4.0]	24.1	93.4	—
q235_div4	Flat	756	29.8 [26.6, 33.1]	[24.6, 35.2]	29.8	0.0	27.1 pp [21.8, 32.4], p=0.000
q235_div4	ACL-Hardened	756	26.7 [23.7, 30.0]	[21.6, 32.1]	29.5	9.4	24.1 pp [19.1, 29.2], p=0.000
q235_div4	AgenticCyOps	756	2.6 [1.7, 4.0]	[1.1, 4.5]	34.8	92.4	—
scout_div4	Flat	378	38.9 [34.1, 43.9]	[31.5, 46.6]	38.9	0.0	38.1 pp [30.2, 45.8], p=0.000
scout_div4	ACL-Hardened	378	33.9 [29.3, 38.8]	[26.5, 41.3]	36.0	5.9	33.1 pp [25.7, 40.5], p=0.000
scout_div4	AgenticCyOps	378	0.8 [0.3, 2.3]	[0.0, 2.4]	38.6	98.0	—

ASR = executed / measurable; attempt = executed + blocked; Δ = paired difference vs AgenticCyOps with cluster-bootstrap CI over variants.

T2a. ASR per attack path, q235_div4, all domains pooled



T2b. ASR per attack path, q235_div4, CyberOps



T2a. Per attack path, q235_div4, all domains

Attack path	Flat ASR %	ACL-Hardened ASR %	AgenticCyOps ASR %	Flat – ACO (pp)	ACL – ACO (pp)
AP-1 Tool redirection	16.7 [9.3, 28.0] (n=60)	0.0 [0.0, 6.0] (n=60)	0.0 [0.0, 6.0] (n=60)	16.7 (p_holm=0.059)	0.0 (p_holm=1.000)
AP-2 Memory poisoning	0.0 [0.0, 7.9] (n=45)	0.0 [0.0, 7.9] (n=45)	0.0 [0.0, 7.9] (n=45)	0.0 (p_holm=1.000)	0.0 (p_holm=1.000)
AP-3 Confused deputy	0.0 [0.0, 6.0] (n=60)	0.0 [0.0, 6.0] (n=60)	0.0 [0.0, 6.0] (n=60)	0.0 (p_holm=1.000)	0.0 (p_holm=1.000)
AP-4 Cross-phase leak	0.0 [0.0, 20.4] (n=15)	0.0 [0.0, 20.4] (n=15)	0.0 [0.0, 20.4] (n=15)	0.0 (p_holm=1.000)	0.0 (p_holm=1.000)
AP-5 Irreversible action	18.3 [10.6, 29.9] (n=60)	18.3 [10.6, 29.9] (n=60)	0.0 [0.0, 6.0] (n=60)	18.3 (p_holm=0.059)	18.3 (p_holm=0.229)
AP-6 Replay	3.3 [0.9, 11.4] (n=60)	0.0 [0.0, 6.0] (n=60)	0.0 [0.0, 6.0] (n=60)	3.3 (p_holm=1.000)	0.0 (p_holm=1.000)
AP-7 Action chain	6.7 [2.6, 15.9] (n=60)	5.0 [1.7, 13.7] (n=60)	6.7 [2.6, 15.9] (n=60)	0.0 (p_holm=1.000)	-1.7 (p_holm=1.000)
AP-8 Parameter manipulation	18.3 [10.6, 29.9] (n=60)	23.3 [14.4, 35.4] (n=60)	5.0 [1.7, 13.7] (n=60)	13.3 (p_holm=0.505)	18.3 (p_holm=0.070)
AP-9 Handoff poisoning	56.7 [44.1, 68.4] (n=60)	41.7 [30.1, 54.3] (n=60)	5.0 [1.7, 13.7] (n=60)	51.7 (p_holm=0.000)	36.7 (p_holm=0.000)
AP-10 Validator manipulation	75.0 [62.8, 84.2] (n=60)	75.0 [62.8, 84.2] (n=60)	3.3 [0.9, 11.4] (n=60)	71.7 (p_holm=0.000)	71.7 (p_holm=0.000)
AP-11 Operational context	20.0 [11.8, 31.8] (n=60)	10.0 [4.7, 20.2] (n=60)	1.7 [0.3, 8.9] (n=60)	18.3 (p_holm=0.016)	8.3 (p_holm=0.889)
AP-12 Concurrent actions	13.3 [6.9, 24.2] (n=60)	16.7 [9.3, 28.0] (n=60)	3.3 [0.9, 11.4] (n=60)	10.0 (p_holm=1.000)	13.3 (p_holm=0.650)
AP-13 Adversarial memory	100.0 [94.0, 100.0] (n=60)	100.0 [94.0, 100.0] (n=60)	5.0 [1.7, 13.7] (n=60)	95.0 (p_holm=0.000)	95.0 (p_holm=0.000)
AP-15 Infrastructure integrity	77.8 [61.9, 88.3] (n=36)	77.8 [61.9, 88.3] (n=36)	5.6 [1.5, 18.1] (n=36)	72.2 (p_holm=0.000)	72.2 (p_holm=0.000)

Wilson 95% intervals; p_holm = Holm-corrected paired p-value within the family of attack paths.


T2b. Per attack path, q235_div4, CyberOps

Attack path	Flat ASR %	ACL-Hardened ASR %	AgenticCyOps ASR %	Flat – ACO (pp)	ACL – ACO (pp)
AP-1 Tool redirection	20.0 [7.0, 45.2] (n=15)	0.0 [0.0, 20.4] (n=15)	0.0 [0.0, 20.4] (n=15)	20.0 (p_holm=1.000)	0.0 (p_holm=1.000)
AP-2 Memory poisoning	0.0 [0.0, 20.4] (n=15)	0.0 [0.0, 20.4] (n=15)	0.0 [0.0, 20.4] (n=15)	0.0 (p_holm=1.000)	0.0 (p_holm=1.000)
AP-3 Confused deputy	0.0 [0.0, 20.4] (n=15)	0.0 [0.0, 20.4] (n=15)	0.0 [0.0, 20.4] (n=15)	0.0 (p_holm=1.000)	0.0 (p_holm=1.000)
AP-5 Irreversible action	6.7 [1.2, 29.8] (n=15)	13.3 [3.7, 37.9] (n=15)	0.0 [0.0, 20.4] (n=15)	6.7 (p_holm=1.000)	13.3 (p_holm=1.000)
AP-6 Replay	13.3 [3.7, 37.9] (n=15)	0.0 [0.0, 20.4] (n=15)	0.0 [0.0, 20.4] (n=15)	13.3 (p_holm=1.000)	0.0 (p_holm=1.000)
AP-7 Action chain	20.0 [7.0, 45.2] (n=15)	6.7 [1.2, 29.8] (n=15)	13.3 [3.7, 37.9] (n=15)	6.7 (p_holm=1.000)	-6.7 (p_holm=1.000)
AP-8 Parameter manipulation	33.3 [15.2, 58.3] (n=15)	46.7 [24.8, 69.9] (n=15)	0.0 [0.0, 20.4] (n=15)	33.3 (p_holm=1.000)	46.7 (p_holm=0.190)
AP-9 Handoff poisoning	46.7 [24.8, 69.9] (n=15)	0.0 [0.0, 20.4] (n=15)	0.0 [0.0, 20.4] (n=15)	46.7 (p_holm=0.190)	0.0 (p_holm=1.000)
AP-10 Validator manipulation	100.0 [79.6, 100.0] (n=15)	100.0 [79.6, 100.0] (n=15)	13.3 [3.7, 37.9] (n=15)	86.7 (p_holm=0.000)	86.7 (p_holm=0.000)
AP-11 Operational context	20.0 [7.0, 45.2] (n=15)	13.3 [3.7, 37.9] (n=15)	0.0 [0.0, 20.4] (n=15)	20.0 (p_holm=1.000)	13.3 (p_holm=1.000)
AP-12 Concurrent actions	26.7 [10.9, 51.9] (n=15)	13.3 [3.7, 37.9] (n=15)	0.0 [0.0, 20.4] (n=15)	26.7 (p_holm=1.000)	13.3 (p_holm=1.000)
AP-13 Adversarial memory	100.0 [79.6, 100.0] (n=15)	100.0 [79.6, 100.0] (n=15)	20.0 [7.0, 45.2] (n=15)	80.0 (p_holm=0.004)	80.0 (p_holm=0.004)
AP-15 Infrastructure integrity	77.8 [45.3, 93.7] (n=9)	77.8 [45.3, 93.7] (n=9)	22.2 [6.3, 54.7] (n=9)	55.6 (p_holm=0.000)	55.6 (p_holm=0.000)


Wilson 95% intervals; p_holm = Holm-corrected paired p-value within the family of attack paths.

T3. Benign utility (E1)

Benign task completed



Benign incidents with any denial



T3. Benign utility (E1)


Group	Config	N	Task completed % [95%]	Any denial % [95%]	Denials / incident	Latency s median / p95	Tokens primary / validator
llama8b_div4	ACL-Hardened	75	100.0 [95.1, 100.0]	6.7 [2.9, 14.7]	0.21	8.2 / 22.9	17326 / 0
llama8b_div4	AgenticCyOps	75	98.7 [92.8, 99.8]	60.0 [48.7, 70.3]	1.61	36.2 / 112.9	8360 / 8491
llama8b_div4	Flat	75	100.0 [95.1, 100.0]	0.0 [0.0, 4.9]	0.00	8.1 / 58.5	17391 / 0
mistral_div3p	ACL-Hardened	75	86.7 [77.2, 92.6]	24.0 [15.8, 34.8]	0.71	32.8 / 61.2	18389 / 0
mistral_div3p	AgenticCyOps	75	78.7 [68.1, 86.4]	62.7 [51.4, 72.7]	1.43	53.4 / 103.6	11821 / 3475
mistral_div3p	Flat	75	80.0 [69.6, 87.5]	0.0 [0.0, 4.9]	0.00	32.2 / 55.1	18758 / 0
q235_div4	ACL-Hardened	105	58.1 [48.5, 67.1]	94.3 [88.1, 97.4]	7.59	34.9 / 57.5	18376 / 0
q235_div4	AgenticCyOps	105	97.1 [91.9, 99.0]	88.6 [81.1, 93.3]	3.40	112.7 / 172.5	11928 / 15895
q235_div4	Flat	105	100.0 [96.5, 100.0]	0.0 [0.0, 3.5]	0.00	36.2 / 67.5	19768 / 0
q235_div4	LLM-judge only	60	95.0 [86.3, 98.3]	81.7 [70.1, 89.4]	1.82	167.3 / 194.8	11937 / 26321
q235_div4	Symbolic only (no L6)	60	25.0 [15.8, 37.2]	96.7 [88.6, 99.1]	7.42	55.1 / 72.3	12184 / 0
scout_div4	ACL-Hardened	75	70.7 [59.6, 79.8]	52.0 [40.9, 62.9]	1.79	51.0 / 64.5	21316 / 0
scout_div4	AgenticCyOps	75	62.7 [51.4, 72.7]	94.7 [87.1, 97.9]	3.44	94.5 / 150.0	12162 / 8651
scout_div4	Flat	75	84.0 [74.1, 90.6]	0.0 [0.0, 4.9]	0.00	51.9 / 64.8	21685 / 0

T3b. Persistent-state sequence (E1b), AgenticCyOps, q235_div4


Domain	State mode	N	Any denial % [95%]	Denials / incident	Task completed % [95%]	Any denial %: first half / second half of sequence
cyberops	isolated	20	70.0 [48.1, 85.5]	1.45	90.0 [69.9, 97.2]	70 / 70
cyberops	persistent	20	100.0 [83.9, 100.0]	7.35	5.0 [0.9, 23.6]	100 / 100
healthcare	isolated	5	100.0 [56.6, 100.0]	7.60	100.0 [56.6, 100.0]	100 / 100
healthcare	persistent	5	100.0 [56.6, 100.0]	11.40	40.0 [11.8, 76.9]	100 / 100
finance	isolated	5	100.0 [56.6, 100.0]	4.20	100.0 [56.6, 100.0]	100 / 100
finance	persistent	5	100.0 [56.6, 100.0]	8.20	40.0 [11.8, 76.9]	100 / 100
legal	isolated	5	100.0 [56.6, 100.0]	4.80	100.0 [56.6, 100.0]	100 / 100
legal	persistent	5	100.0 [56.6, 100.0]	8.20	40.0 [11.8, 76.9]	100 / 100

T4. Ablations (E3): removing one principle at a time

Attack success rate (%)




Benign incidents with any denial (%)



T4. Ablations (E3), q235_div4

Ablation config	Attack trials	ASR % [95%]	Benign trials	Benign any-denial % [95%]
AgenticCyOps	756	2.6 [1.7, 4.1]	105	88.6 [81.1, 93.3]
agenticcyops -P1	189	5.8 [3.3, 10.1]	60	76.7 [64.6, 85.6]
agenticcyops -P2	189	10.6 [7.0, 15.8]	60	73.3 [61.0, 82.9]
agenticcyops -P3	189	18.0 [13.2, 24.1]	60	65.0 [52.4, 75.8]
agenticcyops -P4	189	10.6 [7.0, 15.8]	60	83.3 [72.0, 90.7]
agenticcyops -P5	189	3.7 [1.8, 7.4]	60	75.0 [62.8, 84.2]
LLM-judge only	189	22.2 [16.9, 28.7]	60	81.7 [70.1, 89.4]
Symbolic only (no L6)	189	1.6 [0.5, 4.6]	60	96.7 [88.6, 99.1]

T5. Which layer blocked the attacks (AgenticCyOps)




T5. Which layer blocked the attacks (AgenticCyOps, q235_div4)


Layer	Blocked trials	Share %
P2_target_not_in_evidence	72	29.6
P3_llm_consensus_reject	66	27.2
P4_schema_violation	48	19.8
P2_parameter_rule_violation	20	8.2
P2_critical_asset	5	2.1
P3_handoff_validation	5	2.1
P3_operational_context	5	2.1
P1_config_integrity_violation	4	1.6
P4_metadata_invalid	3	1.2
P4_similarity_reject	3	1.2
P4_write_replay	3	1.2
P2_wildcard_parameter	3	1.2
P3_replay_detection	2	0.8
P1_response_integrity	2	0.8
P3_bulk_action	1	0.4
P3_execution_verification	1	0.4

T6. ASB paired panel replay (same primary outputs, different validator panels)

ASB defended ASR by validator panel



Defended ASR (%) by attack subtype; LLM ASR without defense in brackets




T6. ASB paired panel replay (E6)


Run (group[panel])	Attack subtype	Config	N	LLM ASR %	Defended ASR %
q235_div4_div3	context_manipulation	agenticcyops	150	39.33	0.0
q235_div4_div3	escape_characters	agenticcyops	300	13.67	0.0
q235_div4_div3	fake_completion	agenticcyops	300	39.67	0.0
q235_div4_div3	naive	agenticcyops	450	35.33	1.56
q235_div4_div3	trigger_phrase	agenticcyops	75	36.0	6.67
q235_div4_div4	context_manipulation	agenticcyops	150	39.33	0.0
q235_div4_div4	escape_characters	agenticcyops	300	13.67	0.0
q235_div4_div4	fake_completion	agenticcyops	300	39.67	0.0
q235_div4_div4	naive	agenticcyops	450	35.33	0.0
q235_div4_div4	trigger_phrase	agenticcyops	75	36.0	0.0
q235_div4_drift	context_manipulation	agenticcyops	10	30.0	0.0
q235_div4_drift	context_manipulation	flat	10	30.0	30.0
q235_div4_drift	escape_characters	agenticcyops	10	0.0	0.0
q235_div4_drift	escape_characters	flat	10	20.0	20.0
q235_div4_drift	fake_completion	agenticcyops	10	30.0	0.0
q235_div4_drift	fake_completion	flat	10	40.0	40.0
q235_div4_drift	naive	agenticcyops	15	26.67	6.67
q235_div4_drift	naive	flat	15	26.67	26.67
q235_div4_drift	trigger_phrase	agenticcyops	5	20.0	0.0
q235_div4_drift	trigger_phrase	flat	5	20.0	20.0
q235_div4_lin3	context_manipulation	agenticcyops	150	39.33	0.0
q235_div4_lin3	escape_characters	agenticcyops	300	13.67	7.33
q235_div4_lin3	fake_completion	agenticcyops	300	39.67	8.67
q235_div4_lin3	naive	agenticcyops	450	35.33	11.11
q235_div4_lin3	trigger_phrase	agenticcyops	75	36.0	13.33
q235_div4_single	context_manipulation	agenticcyops	150	39.33	0.0
q235_div4_single	escape_characters	agenticcyops	300	13.67	1.0
q235_div4_single	fake_completion	agenticcyops	300	39.67	2.33
q235_div4_single	naive	agenticcyops	450	35.33	5.78
q235_div4_single	trigger_phrase	agenticcyops	75	36.0	6.67

T7. Cost (attack trials; measured under concurrent load)

Median latency per incident (s)



Tokens per incident, primary + validators (k)



T7. Cost

Group	Config	N	Primary tokens / trial	Validator tokens / trial	Latency s median / p95
llama8b_div4	ACL-Hardened	378	15620	0	6.8 / 21.5
llama8b_div4	AgenticCyOps	378	7120	8303	34.1 / 153.0
llama8b_div4	Flat	378	15679	0	6.7 / 24.8
mistral_div3p	ACL-Hardened	378	15533	0	18.4 / 53.5
mistral_div3p	AgenticCyOps	378	9353	2288	36.1 / 82.5
mistral_div3p	Flat	378	15685	0	19.7 / 49.2
q235_div4	ACL-Hardened	756	14976	0	36.3 / 76.2
q235_div4	AgenticCyOps	756	8644	9932	88.5 / 179.3
q235_div4	Flat	756	16195	0	37.3 / 81.4
q235_div4	LLM-judge only	189	10759	29289	190.5 / 253.5
q235_div4	Symbolic only (no L6)	189	11178	0	51.9 / 73.8
scout_div4	ACL-Hardened	378	19502	0	47.6 / 68.4
scout_div4	AgenticCyOps	378	11328	6070	77.4 / 135.9
scout_div4	Flat	378	19696	0	47.6 / 69.0

T8. Run provenance (1/2)

Group	Domain	Config	Suffix	Git	Freeze tag	Primary	Quant	Panel	T	State	vLLM
llama8b_div4	cyberops	acl_hardened		98dde989a4	defense-freeze-v2-8-g98dde98	meta-llama/Llama-3.1-8B-Instruct	bf16		0.7	isolated	0.21.0
llama8b_div4	cyberops	agenticcyops		98dde989a4	defense-freeze-v2-8-g98dde98	meta-llama/Llama-3.1-8B-Instruct	bf16	div4	0.7	isolated	0.21.0
llama8b_div4	cyberops	flat		98dde989a4	defense-freeze-v2-8-g98dde98	meta-llama/Llama-3.1-8B-Instruct	bf16		0.7	isolated	0.21.0
llama8b_div4	finance	acl_hardened		98dde989a4	defense-freeze-v2-8-g98dde98	meta-llama/Llama-3.1-8B-Instruct	bf16		0.7	isolated	0.21.0
llama8b_div4	finance	agenticcyops		98dde989a4	defense-freeze-v2-8-g98dde98	meta-llama/Llama-3.1-8B-Instruct	bf16	div4	0.7	isolated	0.21.0
llama8b_div4	finance	flat		98dde989a4	defense-freeze-v2-8-g98dde98	meta-llama/Llama-3.1-8B-Instruct	bf16		0.7	isolated	0.21.0
mistral_div3p	cyberops	acl_hardened		98dde989a4	defense-freeze-v2-8-g98dde98	mistralai/Mistral-Small-3.2-24B-Instruct-2506	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
mistral_div3p	cyberops	agenticcyops		98dde989a4	defense-freeze-v2-8-g98dde98	mistralai/Mistral-Small-3.2-24B-Instruct-2506	bf16	mistral_div3p	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
mistral_div3p	cyberops	flat		98dde989a4	defense-freeze-v2-8-g98dde98	mistralai/Mistral-Small-3.2-24B-Instruct-2506	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
mistral_div3p	finance	acl_hardened		98dde989a4	defense-freeze-v2-8-g98dde98	mistralai/Mistral-Small-3.2-24B-Instruct-2506	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
mistral_div3p	finance	agenticcyops		98dde989a4	defense-freeze-v2-8-g98dde98	mistralai/Mistral-Small-3.2-24B-Instruct-2506	bf16	mistral_div3p	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
mistral_div3p	finance	flat		98dde989a4	defense-freeze-v2-8-g98dde98	mistralai/Mistral-Small-3.2-24B-Instruct-2506	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	cyberops	acl_hardened		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	cyberops	agenticcyops		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	cyberops	agenticcyops	_disabled_P1	c66c04b66a	defense-freeze-v2-5-gc66c04b	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	cyberops	agenticcyops	_disabled_P2	c66c04b66a	defense-freeze-v2-5-gc66c04b	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	cyberops	agenticcyops	_disabled_P3	c66c04b66a	defense-freeze-v2-5-gc66c04b	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	cyberops	agenticcyops	_disabled_P4	c66c04b66a	defense-freeze-v2-5-gc66c04b	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	cyberops	agenticcyops	_disabled_P5	c66c04b66a	defense-freeze-v2-5-gc66c04b	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	cyberops	flat		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	cyberops	llm_judge		c66c04b66a	defense-freeze-v2-5-gc66c04b	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	cyberops	symbolic_only		c66c04b66a	defense-freeze-v2-5-gc66c04b	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	finance	acl_hardened		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	finance	agenticcyops		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	finance	flat		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0

T8. Run provenance (2/2)

Group	Domain	Config	Suffix	Git	Freeze tag	Primary	Quant	Panel	T	State	vLLM
q235_div4	healthcare	acl_hardened		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	healthcare	agenticcyops		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	healthcare	flat		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	legal	acl_hardened		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	legal	agenticcyops		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4	legal	flat		8e94f23700	defense-freeze-v2-4-g8e94f23	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
q235_div4_persistent	cyberops	agenticcyops		c66c04b66a	defense-freeze-v2-5-gc66c04b	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	persistent	0.19.1rc1.dev117+g3352bf8b0
q235_div4_persistent	finance	agenticcyops		c66c04b66a	defense-freeze-v2-5-gc66c04b	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	persistent	0.19.1rc1.dev117+g3352bf8b0
q235_div4_persistent	healthcare	agenticcyops		c66c04b66a	defense-freeze-v2-5-gc66c04b	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	persistent	0.19.1rc1.dev117+g3352bf8b0
q235_div4_persistent	legal	agenticcyops		c66c04b66a	defense-freeze-v2-5-gc66c04b	Qwen/Qwen3-235B-A22B-Instruct-2507	bf16	div4	0.7	persistent	0.19.1rc1.dev117+g3352bf8b0
scout_div4	cyberops	acl_hardened		a369d18c36	defense-freeze-v2.1	meta-llama/Llama-4-Scout-17B-16E-Instruct	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
scout_div4	cyberops	agenticcyops		a369d18c36	defense-freeze-v2.1	meta-llama/Llama-4-Scout-17B-16E-Instruct	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
scout_div4	cyberops	flat		a369d18c36	defense-freeze-v2.1	meta-llama/Llama-4-Scout-17B-16E-Instruct	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
scout_div4	finance	acl_hardened		a369d18c36	defense-freeze-v2.1	meta-llama/Llama-4-Scout-17B-16E-Instruct	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
scout_div4	finance	agenticcyops		a369d18c36	defense-freeze-v2.1	meta-llama/Llama-4-Scout-17B-16E-Instruct	bf16	div4	0.7	isolated	0.19.1rc1.dev117+g3352bf8b0
scout_div4	finance	flat		a369d18c36	defense-freeze-v2.1	meta-llama/Llama-4-Scout-17B-16E-Instruct	bf16		0.7	isolated	0.19.1rc1.dev117+g3352bf8b0

Scope and notes

- ASR = executed / measurable trials. Trials scored not_measurable (no canary for the attack's effect) or error are excluded from every rate; attempt = executed + blocked.
- AP-4 and AP-14 have no measurable canary in CyberOps, finance and legal and are reported only where the oracle scores them (healthcare AP-4).
- Every (domain, config) ran as its own stream with its own tool servers, memory gateway and vector store; per-trial state isolation (isolated mode) except the E1b persistent runs.
- Latency and token figures were measured under 12-19 concurrent streams sharing the primary model server, so absolute latencies are load-dependent; ratios between configs within a group are comparable.
- Tool servers return scripted responses; the primary model sees one storyline per incident, which bounds the attempt rate.
- ASB panels are a paired replay of the same primary outputs through each validator panel; the drift check (live vs replayed, 50 cases) is in T6 as *_drift.
- Held-out variants (E4) and the adaptive attacker (E5) were not run.
- Defense code is frozen at defense-freeze-v2.1 (v2 plus an agent-output parser fix; the frozen directories are identical between the two tags).

Per-run reports: results/eval_attacks/group_<G>/<domain>/attack_report.pdf; ASB panel analytics: results/asb/**/asb_analytics.pdf.